732A54 – Big Data Analytics: SparkSQL

Version: Dec 1, 2016

DataFrames

A DataFrame is a distributed collection of data organized into named columns. It is conceptually equivalent to a table in a relational database or a data frame in R/Python, but with richer optimizations under the hood. DataFrames can be constructed from a wide array of sources such as: structured data files, tables in Hive, external databases, or existing RDDs.

http://spark.apache.org/docs/1.6.0/sql-programming-guide.html

http://spark.apache.org/docs/1.6.0/api/python/pyspark.sql.html

LINKÖPING UNIVERSITY

SQLContext & HiveContext

- Start with obtaining `SparkContext` object and then `SQLContext` from it

```
sc = SparkContext()

sqlContext = SQLContext(sc)
```

- HiveContext provides additional features to SQLContext (likely not needed for the lab assignment)

```
from pyspark.sql import HiveContext

sqlContext = HiveContext(sc)
```

LINKÖPING UNIVERSITY

Imports

Don't forget to import relevant classes first!

```
from pyspark import SparkContext
from pyspark.sql import SQLContext, Row
from pyspark.sql import functions as F
```

LINKÖPING
UNIVERSITY

Create a DataFrame from a RDD

- Two ways:
  - Inferring the schema using reflection
  - Specifying the schema programatically

- Then register the table

li.u LINKÖPING
UNIVERSITY

Create a DataFrame from a RDD – way I

```
# Load a text file and convert each line to a Row.

rdd = sc.textFile("FILENAME")

parts = rdd.map(lambda l: l.split(";"))

tempReadings = parts.map(lambda p: Row(station=p[0],
date=p[1], year=p[1].split("-")[0], time=p[2],
value=float(p[3]), quality=p[4] ))
```

- Inferring the schema and registering the DataFrame as a table

```
schemaTempReadings =
sqlContext.createDataFrame(tempReadings)

schemaTempReadings.registerTempTable("tempReadings")


# Can run queries now
```

LiU LINKÖPING
UNIVERSITY

Create a DataFrame from a RDD – way II

```
# Load a text file and convert each line to a tuple.

rdd = sc.textFile("FILENAME")

parts = rdd.map(lambda l: l.split(";"))

tempReadingsRow = parts.map(lambda p: (p[0], p[1], int(p[1].split("-")[0]),
int(p[1].split("-")[1]), p[2], float(p[3]), p[4] ))
```

- Specifying the schema programatically and registering the DataFrame as a table

```
tempReadingsString = ["station", "date", "year", "month", "time", "value",
"quality"]

# Apply the schema to the RDD.

schemaTempReadings = sqlContext.createDataFrame(tempReadingsRow,
tempReadingsString)

# Register the DataFrame as a table.

schemaTempReadings.registerTempTable("tempReadingsTable")

# Can run queries now
```

Access columns' values

- In Python it's possible to access a DataFrame's columns either by attribute (`dataframe.age`) or by indexing (`dataframe['age']`). While the former is convenient for interactive data exploration, users are highly encouraged to **use the latter form**, which is future proof and won't break with column names that are also attributes on the DataFrame class.

Run SQL queries

- Two ways:

  - Write regular SQL over DataFrames registered as tables using `sqlContext.sql("Your regular SQL query")`

  - Use API methods, such as `select()`, `filter()`, `groupBy()`, `agg()`, etc.

The results of SQL queries are RDDs and support all the normal RDD operations.

```
max1950rdd = max1950.map(lambda p: "Name: " + p.name)

for value in max1950rdd.collect():

  print(value)
```

Run SQL queries – regular SQL – way I

These examples use **temperature-readings.csv**

\# SQL can be run over DataFrames that have been registered as a table.

max1950 = sqlContext.sql("SELECT max(value) as value FROM tempReadings WHERE year=1950")

largerThan10Degrees = sqlContext.sql("SELECT year, month, count(value) as value FROM tempReadingsTable WHERE year=1950 and value>=10.0 group by year, month")

Run SQL queries – API methods – way II

These examples use **temperature-readings.csv**

```
schemaTempReadingsMin =
schemaTempReadings.groupBy('year', 'month', 'day',
'station').agg(F.min('value').alias('dailymin'))
.orderBy(['year', 'month', 'day', 'station'],
ascending=[0,0,0,1])
```

Run SQL queries – API methods – way II

This example uses **temperature-readings.csv and stations-Ostergotland.csv**

Note: Partial code !!!!! - it demonstrates the join not the entire code needed before

```
# Define the schema

stationsString = ["stationNum", "stationName"]

precReadingsString = ["station", "date", "year", "month", "day",
"time", "value", "quality"]

# Apply the schema to the RDD.

schemaStations = sqlContext.createDataFrame(stationsRow,
stationsString)

schemaPrecReadings = sqlContext.createDataFrame(precReadingsRow,
precReadingsString)

precStationOst = schemaStations.join(schemaPrecReadings,
schemaStations['stationNum']==schemaPrecReadings['station'], 'inner')
```